

Implement K-Means clustering/ hierarchical clustering on sales_data_sample.csv dataset. Determine the number of clusters using the elbow method.

In [198]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
#Importing the required libraries.
```

In [199]:

```
from sklearn.cluster import KMeans, k_means #For clustering
from sklearn.decomposition import PCA #Linear Dimensionality reduction.
```

In [200]:

```
df = pd.read_csv("sales_data_sample.csv") #Loading the dataset.
```

Preprocessing

In [201]:

```
df.head()
```

Out[201]:

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE	STATUS	QTR_ID	MONTH
0	10107	30	95.70	2	2871.00	2/24/2003 0:00	Shipped	1	
1	10121	34	81.35	5	2765.90	5/7/2003 0:00	Shipped	2	
2	10134	41	94.74	2	3884.34	7/1/2003 0:00	Shipped	3	
3	10145	45	83.26	6	3746.70	8/25/2003 0:00	Shipped	3	
4	10159	49	100.00	14	5205.27	10/10/2003 0:00	Shipped	4	

5 rows x 25 columns

In [202]:

```
df.shape
```

Out[202]:

(2823, 25)


In [203]:

```
df.describe()
```

Out[203]:

ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID
-------------	-----------------	-----------	-----------------	-------	--------	----------

count	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	2823.SALES	2823.QTR_ID	2823.MONTH_ID	28
mean	10258.725115	35.092809	83.658544	6.466171	3553.889072	2.717676	7.092455	20
std	92.085478	9.741443	20.174277	4.225841	1841.865106	1.203878	3.656633	
min	10100.000000	6.000000	26.880000	1.000000	482.130000	1.000000	1.000000	20
25%	10180.000000	27.000000	68.860000	3.000000	2203.430000	2.000000	4.000000	20
50%	10262.000000	35.000000	95.700000	6.000000	3184.800000	3.000000	8.000000	20
75%	10333.500000	43.000000	100.000000	9.000000	4508.000000	4.000000	11.000000	20
max	10425.000000	97.000000	100.000000	18.000000	14082.800000	4.000000	12.000000	20



In [204]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2823 entries, 0 to 2822
Data columns (total 25 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ORDERNUMBER         2823 non-null   int64
1   QUANTITYORDERED     2823 non-null   int64
2   PRICEEACH           2823 non-null   float64
3   ORDERLINENUMBER     2823 non-null   int64
4   SALES                2823 non-null   float64
5   ORDERDATE           2823 non-null   object
6   STATUS              2823 non-null   object
7   QTR_ID              2823 non-null   int64
8   MONTH_ID            2823 non-null   int64
9   YEAR_ID             2823 non-null   int64
10  PRODUCTLINE         2823 non-null   object
11  MSRP                2823 non-null   int64
12  PRODUCTCODE         2823 non-null   object
13  CUSTOMERNAME        2823 non-null   object
14  PHONE               2823 non-null   object
15  ADDRESSLINE1        2823 non-null   object
16  ADDRESSLINE2        302 non-null    object
17  CITY                2823 non-null   object
18  STATE               1337 non-null   object
19  POSTALCODE          2747 non-null   object
20  COUNTRY             2823 non-null   object
21  TERRITORY           1749 non-null   object
22  CONTACTLASTNAME     2823 non-null   object
23  CONTACTFIRSTNAME    2823 non-null   object
24  DEALSIZE            2823 non-null   object
dtypes: float64(2), int64(7), object(16)
memory usage: 551.5+ KB
```

In [205]:

```
df.isnull().sum()
```

```
Out[205]:
ORDERNUMBER         0
QUANTITYORDERED     0
PRICEEACH           0
ORDERLINENUMBER     0
SALES               0
ORDERDATE           0
STATUS              0
QTR_ID             0
MONTH_ID            0
YEAR_ID            0
PRODUCTLINE         0
MSRP               0
PRODUCTCODE         0
CUSTOMERNAME        0
PHONE              0
```

```
PHONE 0
ADDRESSLINE1 0
ADDRESSLINE2 2521
CITY 0
STATE 1486
POSTALCODE 76
COUNTRY 0
TERRITORY 1074
CONTACTLASTNAME 0
CONTACTFIRSTNAME 0
DEALSIZE 0
dtype: int64
```

In [206]:

```
df.dtypes
```

Out[206]:

```
ORDERNUMBER          int64
QUANTITYORDERED      int64
PRICEEACH            float64
ORDERLINENUMBER      int64
SALES                float64
ORDERDATE            object
STATUS              object
QTR_ID              int64
MONTH_ID            int64
YEAR_ID             int64
PRODUCTLINE         object
MSRP                int64
PRODUCTCODE         object
CUSTOMERNAME        object
PHONE              object
ADDRESSLINE1        object
ADDRESSLINE2        object
CITY                object
STATE              object
POSTALCODE          object
COUNTRY            object
TERRITORY          object
CONTACTLASTNAME     object
CONTACTFIRSTNAME    object
DEALSIZE            object
dtype: object
```

In [207]:

```
df_drop = ['ADDRESSLINE1', 'ADDRESSLINE2', 'STATUS', 'POSTALCODE', 'CITY', 'TERRITORY',
'PHONE', 'STATE', 'CONTACTFIRSTNAME', 'CONTACTLASTNAME', 'CUSTOMERNAME', 'ORDERNUMBER']
df = df.drop(df_drop, axis=1) #Dropping the categorical unecessary columns along with c
olumns having null values. Can't fill the null values are there are alot of null values.
```

In [208]:

```
df.isnull().sum()
```

Out[208]:

```
QUANTITYORDERED      0
PRICEEACH            0
ORDERLINENUMBER      0
SALES                0
ORDERDATE            0
QTR_ID              0
MONTH_ID            0
YEAR_ID             0
PRODUCTLINE         0
MSRP                0
PRODUCTCODE         0
COUNTRY            0
DEALSIZE            0
dtype: int64
```

In [209]:

```
df.dtypes
```

Out[209]:

```
QUANTITYORDERED    int64
PRICEEACH           float64
ORDERLINENUMBER     int64
SALES               float64
ORDERDATE           object
QTR_ID              int64
MONTH_ID            int64
YEAR_ID             int64
PRODUCTLINE         object
MSRP                int64
PRODUCTCODE         object
COUNTRY             object
DEALSIZE            object
dtype: object
```

In []:

```
# Checking the categorical columns.
```

In [210]:

```
df['COUNTRY'].unique()
```

Out[210]:

```
array(['USA', 'France', 'Norway', 'Australia', 'Finland', 'Austria', 'UK',
       'Spain', 'Sweden', 'Singapore', 'Canada', 'Japan', 'Italy',
       'Denmark', 'Belgium', 'Philippines', 'Germany', 'Switzerland',
       'Ireland'], dtype=object)
```

In [211]:

```
df['PRODUCTLINE'].unique()
```

Out[211]:

```
array(['Motorcycles', 'Classic Cars', 'Trucks and Buses', 'Vintage Cars',
       'Planes', 'Ships', 'Trains'], dtype=object)
```

In [212]:

```
df['DEALSIZE'].unique()
```

Out[212]:

```
array(['Small', 'Medium', 'Large'], dtype=object)
```

In [213]:

```
productline = pd.get_dummies(df['PRODUCTLINE']) #Converting the categorical columns.
Dealsize = pd.get_dummies(df['DEALSIZE'])
```

In [214]:

```
df = pd.concat([df,productline,Dealsize], axis = 1)
```

In [215]:

```
df_drop = ['COUNTRY','PRODUCTLINE','DEALSIZE'] #Dropping Country too as there are alot o
f countries.
df = df.drop(df_drop, axis=1)
```

In [216]:

```
df['PRODUCTCODE'] = pd.Categorical(df['PRODUCTCODE']).codes #Converting the datatype.
```

In [217]:

```
df.drop('ORDERDATE', axis=1, inplace=True) #Dropping the Orderdate as Month is already in  
cluded.
```

In [218]:

```
df.dtypes #All the datatypes are converted into numeric
```

Out[218]:

```
QUANTITYORDERED    int64
PRICEEACH           float64
ORDERLINENUMBER     int64
SALES               float64
QTR_ID             int64
MONTH_ID            int64
YEAR_ID            int64
MSRP               int64
PRODUCTCODE         int8
Classic Cars        uint8
Motorcycles          uint8
Planes              uint8
Ships               uint8
Trains              uint8
Trucks and Buses    uint8
Vintage Cars        uint8
Large               uint8
Medium              uint8
Small               uint8
dtype: object
```


Plotting the Elbow Plot to determine the number of clusters.


In [219]:

```
distortions = [] # Within Cluster Sum of Squares from the centroid
K = range(1,10)
for k in K:
    kmeanModel = KMeans(n_clusters=k)
    kmeanModel.fit(df)
    distortions.append(kmeanModel.inertia_) #Appending the inertia to the Distortions
```

In [220]:

```
plt.figure(figsize=(16,8))
plt.plot(K, distortions, 'bx-')
plt.xlabel('k')
plt.ylabel('Distortion')
plt.title('The Elbow Method showing the optimal k')
plt.show()
```





As the number of k increases Inertia decreases.

Observations: A Elbow can be observed at 3 and after that the curve decreases gradually.

In [221]:

```
X_train = df.values #Returns a numpy array.
```

In [222]:

```
X_train.shape
```

Out[222]:

```
(2823, 19)
```

In [223]:

```
model = KMeans(n_clusters=3,random_state=2) #Number of cluster = 3
model = model.fit(X_train) #Fitting the values to create a model.
predictions = model.predict(X_train) #Predicting the cluster values (0,1,or 2)
```

In [225]:

```
unique,counts = np.unique(predictions,return_counts=True)
```

In [226]:

```
counts = counts.reshape(1,3)
```

In [227]:

```
counts_df = pd.DataFrame(counts,columns=['Cluster1','Cluster2','Cluster3'])
```

In [228]:

```
counts_df.head()
```

Out[228]:

	Cluster1	Cluster2	Cluster3
0	1083	1367	373

Visualization

In [229]:

```
pca = PCA(n_components=2) #Converting all the features into 2 columns to make it easy to visualize using Principal Component Analysis.
```

In [230]:


```
reduced_X = pd.DataFrame(pca.fit_transform(X_train),columns=['PCA1','PCA2']) #Creating a DataFrame.
```

In [231]:

Out [231] :

In [232]:

```
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'], reduced_X['PCA2'])
```



In [233]:

```
[ 3.08302853e+01,  7.00755230e+01,  6.67300658e+00,
 2.12409474e+03,  2.71762985e+00,  7.09509876e+00,
 2.00381127e+03,  7.84784199e+01,  6.24871982e+01,
 2.64813460e-01,  1.21433797e-01,  1.29480614e-01,
 1.00219459e-01,  3.87710315e-02,  9.21726408e-02,
 2.53108998e-01,  6.93889390e-18,  6.21799561e-02,
 9.37820044e-01],
 [ 4.45871314e+01,  9.98931099e+01,  5.75603217e+00,
 7.09596863e+03,  2.71045576e+00,  7.06434316e+00,
 2.00389008e+03,  1.45823056e+02,  3.14959786e+01,
 5.33512064e-01,  1.07238606e-01,  7.23860590e-02,
 2.14477212e-02,  1.07238606e-02,  1.31367292e-01,
 1.23324397e-01,  4.20911528e-01,  5.79088472e-01,
 5.55111512e-17]])
```

In [234]:

```
reduced_centers = pca.transform(model.cluster_centers_) #Transforming the centroids into
3 in x and y coordinates
```

In [235]:

```
reduced_centers
```

Out[235]:


```
array([[ 5.84994044e+02, -4.36786931e+00],
       [-1.43005891e+03,  2.60041009e+00],
       [ 3.54247180e+03,  3.15185487e+00]])
```

In [236]:

```
plt.figure(figsize=(14,10))
plt.scatter(reduced_X['PCA1'],reduced_X['PCA2'])
plt.scatter(reduced_centers[:,0],reduced_centers[:,1],color='black',marker='x',s=300) #Plotting the centroids
```

Out[236]:

```
<matplotlib.collections.PathCollection at 0x218deb6e220>
```



In [237]:

```
reduced_X['Clusters'] = predictions #Adding the Clusters to the reduced dataframe.
```

In [238]:

```
reduced_X.head()
```

Out[238]:

	PCA1	PCA2	Clusters
0	-682.488323	-42.819535	1
1	-787.665502	-41.694991	1
2	330.732170	-26.481208	0
3	193.040232	-26.285766	0
4	1651.532874	-6.891196	0


In [239]:


```
#Plotting the clusters
plt.figure(figsize=(14,10))
# taking the cluster number and first column taking the sa
me cluster number and second column Assigning the color
plt.scatter(reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 0].loc[:, 'PCA2'], color='slateblue')
plt.scatter(reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 1].loc[:, 'PCA2'], color='springgreen')
plt.scatter(reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA1'], reduced_X[reduced_X['Clusters'] == 2].loc[:, 'PCA2'], color='indigo')

plt.scatter(reduced_centers[:,0], reduced_centers[:,1], color='black', marker='x', s=300)
```

Out[239]:

<matplotlib.collections.PathCollection at 0x218dce9e1f0>





In []: