

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months.

Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling> Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix.

In [46]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt #Importing the libraries
```

In [47]:

```
df = pd.read_csv("Churn_Modelling.csv")
```

Preprocessing.

In [48]:

```
df.head()
```

Out[48]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125510.82	1	

In [49]:

```
df.shape
```

Out[49]:

```
(10000, 14)
```

In [50]:

```
df.describe()
```

Out[50]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsChurn
0	15634602	15634602	619	42	2	0.00	1	1	0

count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.00000	Is
RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	Is		
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550		
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584		
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.00000		
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.00000		
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.00000		
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.00000		
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.00000		



In [51]:

```
df.isnull()
```

Out[51]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrC
0	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False
...
9995	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns



In [52]:

```
df.isnull().sum()
```

Out[52]:

RowNumber	0
CustomerId	0
Surname	0
CreditScore	0
Geography	0
Gender	0
Age	0
Tenure	0
Balance	0
NumOfProducts	0
HasCrCard	0
IsActiveMember	0
EstimatedSalary	0
Exited	0
dtype:	int64

In [53]:

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):

```

# Column Non-Null Count Dtype
---
0 RowNumber 10000 non-null int64
1 CustomerId 10000 non-null int64
2 Surname 10000 non-null object
3 CreditScore 10000 non-null int64
4 Geography 10000 non-null object
5 Gender 10000 non-null object
6 Age 10000 non-null int64
7 Tenure 10000 non-null int64
8 Balance 10000 non-null float64
9 NumOfProducts 10000 non-null int64
10 HasCrCard 10000 non-null int64
11 IsActiveMember 10000 non-null int64
12 EstimatedSalary 10000 non-null float64
13 Exited 10000 non-null int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB

```

In [54]:

```
df.dtypes
```

Out[54]:

```

RowNumber      int64
CustomerId      int64
Surname        object
CreditScore     int64
Geography       object
Gender          object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts  int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object

```

In [55]:

```
df.columns
```

Out[55]:

```

Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')

```

In [56]:

```
df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis= 1) #Dropping the unnecessary columns
```

In [57]:

```
df.head()
```

Out[57]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	France	Female	42	2	0.00	1	1	1	101348.88
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	699	France	Female	39	1	0.00	2	0	0	93826.63
4	950	Spain	Female	42	2	125510.92	1	1	1	70094.10

4	630	Spain	Female	43	2	123510.62	1	1	79004.10	
	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary

Visualization

In [101]:


```
def visualization(x, y, xlabel):
    plt.figure(figsize=(10,5))
    plt.hist([x, y], color=['red', 'green'], label = ['exit', 'not_exit'])
    plt.xlabel(xlabel, fontsize=20)
    plt.ylabel("No. of customers", fontsize=20)
    plt.legend()
```

In [102]:

```
df_churn_exited = df[df['Exited']==1]['Tenure']
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

In [103]:

```
visualization(df_churn_exited, df_churn_not_exited, "Tenure")
```




In [105]:

```
df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

In [106]:

```
visualization(df_churn_exited2, df_churn_not_exited2, "Age")
```



Converting the Categorical Variables

In [59]:

```
X = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
states = pd.get_dummies(df[['Geography']], drop_first = True)
gender = pd.get_dummies(df[['Gender']], drop_first = True)
```

In [61]:

```
df = pd.concat([df, gender, states], axis = 1)
```

Splitting the training and testing Dataset

In [62]:

```
df.head()
```

Out[62]:

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary
0	619	France	Female	42	2	0.00	1	1	1	101348.88
1	608	Spain	Female	41	1	83807.86	1	0	1	112542.58
2	502	France	Female	42	8	159660.80	3	1	0	113931.57
3	699	France	Female	39	1	0.00	2	0	0	93826.63
4	850	Spain	Female	43	2	125510.82	1	1	1	79084.10

In [63]:

```
X = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Male', 'Germany', 'Spain']]
```

In [64]:

```
y = df[['Exited']]
```

In [65]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30)
```

Normalizing the values with mean as 0 and Standard Deviation as 1

In [66]:

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

In [67]:

```
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

In [68]:

```
X_train
```

```
Out[68]:
```

```
array([[ 4.56838557e-01, -9.45594735e-01,  1.58341939e-03, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
       [-2.07591864e-02, -2.77416637e-01,  3.47956411e-01, ...,
        -1.09507222e+00, -5.81969145e-01,  1.74334114e+00],
       [-1.66115021e-01,  1.82257167e+00, -1.38390855e+00, ...,
        -1.09507222e+00, -5.81969145e-01, -5.73611200e-01],
       ...,
       [-3.63383654e-01, -4.68324665e-01,  1.73344838e+00, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
       [ 4.67221117e-01, -1.42286480e+00,  1.38707539e+00, ...,
         9.13181783e-01, -5.81969145e-01,  1.74334114e+00],
       [-8.82511636e-01,  2.95307447e-01, -6.91162564e-01, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01]])
```

```
In [69]:
```

```
X_test
```

```
Out[69]:
```

```
array([[ 3.63395520e-01,  1.99853433e-01,  1.58341939e-03, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
       [-4.15243057e-02,  4.86215475e-01,  1.58341939e-03, ...,
        -1.09507222e+00, -5.81969145e-01,  1.74334114e+00],
       [-1.87923736e+00, -3.72870651e-01, -1.38390855e+00, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
       ...,
       [-6.02182526e-01, -5.63778679e-01, -1.73028154e+00, ...,
        -1.09507222e+00, -5.81969145e-01, -5.73611200e-01],
       [ 1.51585964e+00, -6.59232693e-01,  1.73344838e+00, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01],
       [-5.19122049e-01,  1.04399419e-01,  1.73344838e+00, ...,
         9.13181783e-01, -5.81969145e-01, -5.73611200e-01]])
```

Building the Classifier Model using Keras

```
In [70]:
```

```
import keras #Keras is the wrapper on the top of tensorflow
#Can use Tensorflow as well but won't be able to understand the errors initially.
```

```
In [71]:
```

```
from keras.models import Sequential #To create sequential neural network
from keras.layers import Dense #To create hidden layers
```

```
In [72]:
```

```
classifier = Sequential()
```

```
In [74]:
```

```
#To add the layers
#Dense helps to construct the neurons
#Input Dimension means we have 11 features
# Units is to create the hidden layers
#Uniform helps to distribute the weight uniformly
classifier.add(Dense(activation = "relu",input_dim = 11,units = 6,kernel_initializer = "
uniform"))
```

```
In [75]:
```

```
classifier.add(Dense(activation = "relu",units = 6,kernel_initializer = "uniform")) #A
dding second hidden layers
```

In [76]:

```
classifier.add(Dense(activation = "sigmoid",units = 1,kernel_initializer = "uniform")) #  
Final neuron will be having sigmoid function
```

In [77]:

```
classifier.compile(optimizer="adam",loss = 'binary_crossentropy',metrics = ['accuracy'])  
#To compile the Artificial Neural Network. Used Binary crossentropy as we just have only  
two output
```

In [79]:

```
classifier.summary() #3 layers created. 6 neurons in 1st,6neurons in 2nd layer and 1 neur  
on in last
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_3 (Dense)	(None, 6)	72
dense_4 (Dense)	(None, 6)	42
dense_5 (Dense)	(None, 1)	7
=====	=====	=====
Total params: 121		
Trainable params: 121		
Non-trainable params: 0		
=====		

In [89]:

```
classifier.fit(X_train,y_train,batch_size=10,epochs=50) #Fitting the ANN to training data  
set
```

```
Epoch 1/50  
700/700 [=====] - 0s 674us/step - loss: 0.4293 - accuracy: 0.794  
7  
Epoch 2/50  
700/700 [=====] - 0s 647us/step - loss: 0.4239 - accuracy: 0.794  
7  
Epoch 3/50  
700/700 [=====] - 0s 657us/step - loss: 0.4203 - accuracy: 0.806  
7  
Epoch 4/50  
700/700 [=====] - 0s 664us/step - loss: 0.4167 - accuracy: 0.826  
0  
Epoch 5/50  
700/700 [=====] - 0s 674us/step - loss: 0.4153 - accuracy: 0.828  
7  
Epoch 6/50  
700/700 [=====] - 0s 653us/step - loss: 0.4137 - accuracy: 0.831  
0  
Epoch 7/50  
700/700 [=====] - 0s 658us/step - loss: 0.4125 - accuracy: 0.831  
7  
Epoch 8/50  
700/700 [=====] - 1s 842us/step - loss: 0.4116 - accuracy: 0.830  
6  
Epoch 9/50  
700/700 [=====] - 0s 671us/step - loss: 0.4103 - accuracy: 0.833  
1  
Epoch 10/50  
700/700 [=====] - 0s 682us/step - loss: 0.4100 - accuracy: 0.832  
6  
Epoch 11/50  
700/700 [=====] - 0s 690us/step - loss: 0.4093 - accuracy: 0.833  
7  
Epoch 12/50  
700/700 [=====] - 0s 688us/step - loss: 0.4087 - accuracy: 0.833  
9
```

Epoch 13/50
700/700 [=====] - 0s 675us/step - loss: 0.4081 - accuracy: 0.834
1

Epoch 14/50
700/700 [=====] - 1s 722us/step - loss: 0.4071 - accuracy: 0.833
1

Epoch 15/50
700/700 [=====] - 1s 811us/step - loss: 0.4065 - accuracy: 0.834
1

Epoch 16/50
700/700 [=====] - 0s 711us/step - loss: 0.4056 - accuracy: 0.835
6

Epoch 17/50
700/700 [=====] - 0s 702us/step - loss: 0.4046 - accuracy: 0.836
6

Epoch 18/50
700/700 [=====] - 0s 688us/step - loss: 0.4035 - accuracy: 0.834
3

Epoch 19/50
700/700 [=====] - 1s 715us/step - loss: 0.4024 - accuracy: 0.836
3

Epoch 20/50
700/700 [=====] - 0s 714us/step - loss: 0.4020 - accuracy: 0.833
7

Epoch 21/50
700/700 [=====] - 0s 705us/step - loss: 0.4010 - accuracy: 0.837
4

Epoch 22/50
700/700 [=====] - 1s 720us/step - loss: 0.4003 - accuracy: 0.837
0

Epoch 23/50
700/700 [=====] - 0s 692us/step - loss: 0.3993 - accuracy: 0.837
4

Epoch 24/50
700/700 [=====] - 0s 709us/step - loss: 0.3990 - accuracy: 0.835
6

Epoch 25/50
700/700 [=====] - 1s 871us/step - loss: 0.3984 - accuracy: 0.836
6

Epoch 26/50
700/700 [=====] - 1s 719us/step - loss: 0.3984 - accuracy: 0.836
7

Epoch 27/50
700/700 [=====] - 1s 719us/step - loss: 0.3980 - accuracy: 0.836
6

Epoch 28/50
700/700 [=====] - 0s 695us/step - loss: 0.3981 - accuracy: 0.836
6

Epoch 29/50
700/700 [=====] - 0s 667us/step - loss: 0.3976 - accuracy: 0.837
4

Epoch 30/50
700/700 [=====] - 0s 669us/step - loss: 0.3972 - accuracy: 0.837
3

Epoch 31/50
700/700 [=====] - 0s 670us/step - loss: 0.3970 - accuracy: 0.837
0

Epoch 32/50
700/700 [=====] - 1s 720us/step - loss: 0.3972 - accuracy: 0.837
6

Epoch 33/50
700/700 [=====] - 0s 675us/step - loss: 0.3965 - accuracy: 0.836
7

Epoch 34/50
700/700 [=====] - 0s 680us/step - loss: 0.3961 - accuracy: 0.836
4

Epoch 35/50
700/700 [=====] - 0s 685us/step - loss: 0.3962 - accuracy: 0.837
9

Epoch 36/50
700/700 [=====] - 1s 771us/step - loss: 0.3960 - accuracy: 0.837
0


```
Epoch 37/50
700/700 [=====] - 1s 1ms/step - loss: 0.3963 - accuracy: 0.8366
Epoch 38/50
700/700 [=====] - 1s 764us/step - loss: 0.3962 - accuracy: 0.837
3
Epoch 39/50
700/700 [=====] - 1s 823us/step - loss: 0.3950 - accuracy: 0.838
4
Epoch 40/50
700/700 [=====] - 1s 759us/step - loss: 0.3956 - accuracy: 0.836
1
Epoch 41/50
700/700 [=====] - 1s 773us/step - loss: 0.3949 - accuracy: 0.836
6
Epoch 42/50
700/700 [=====] - 0s 695us/step - loss: 0.3953 - accuracy: 0.836
9
Epoch 43/50
700/700 [=====] - 0s 701us/step - loss: 0.3952 - accuracy: 0.836
9
Epoch 44/50
700/700 [=====] - 0s 707us/step - loss: 0.3952 - accuracy: 0.836
6
Epoch 45/50
700/700 [=====] - 0s 680us/step - loss: 0.3955 - accuracy: 0.837
6
Epoch 46/50
700/700 [=====] - 0s 665us/step - loss: 0.3947 - accuracy: 0.837
3
Epoch 47/50
700/700 [=====] - 0s 708us/step - loss: 0.3947 - accuracy: 0.837
1
Epoch 48/50
700/700 [=====] - 0s 681us/step - loss: 0.3944 - accuracy: 0.837
1
Epoch 49/50
700/700 [=====] - 0s 678us/step - loss: 0.3947 - accuracy: 0.838
3
Epoch 50/50
700/700 [=====] - 1s 869us/step - loss: 0.3944 - accuracy: 0.837
0
```

Out[89]:

```
<tensorflow.python.keras.callbacks.History at 0x1fb1eb93df0>
```

In [90]:

```
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5) #Predicting the result
```

In [97]:

```
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

In [92]:

```
cm = confusion_matrix(y_test, y_pred)
```

In [93]:

```
cm
```

Out[93]:

```
array([[2328,    72],
       [ 425,   175]], dtype=int64)
```

In [94]:

```
accuracy = accuracy_score(y_test, y_pred)
```

In [95]:

```
accuracy
```

Out[95]:


0.8343333333333334

In [98]:

```
plt.figure(figsize = (10,7))
sns.heatmap(cm,annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[98]:

Text(69.0, 0.5, 'Truth')



In [100]:

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.97	0.90	2400
1	0.71	0.29	0.41	600
accuracy			0.83	3000
macro avg	0.78	0.63	0.66	3000
weighted avg	0.82	0.83	0.81	3000

In []: